

CMSC 12300 Final Project

Andy Liu and Nelson Auner

June 4, 2013

1 Introduction

Our project investigates the use of census data to predict whether a household self-reports that its income is over or under \$50,000. Our goals, put succinctly, were to

1. select a subset of prediction methods to examine, based on interpretability of results and computational feasibility
2. test relevant models and select the best one, and,
3. develop an accurate estimation of our model's prediction of a similar, yet different data set.

To this end, we used R and Python, deployed on personal computers as well as AWS, to implement logistical regression of various combinations of regressors. We took the best models, as defined by their performance within the training set, and subjected them to cross-validation. To deal with missing data, we utilized k nearest neighbors to "fill in" missing variables with the nearest closest observation.

Our results suggest the following: Education as a stand-alone predictor is a very good model, with around 5.4% false positive rate and .5% false negative rate. By very good, however, we mean compared to other non-trivial models, considering that the relative lack of diversity in our responses, given an absolute error weighting, favors a classification scheme with all-negative predictions.

2 Description of the Data set

Our analysis is on the Census-Income (KDD) Data Set, a classic machine learning dataset of census data from the 1994 and 1995 Current Population Surveys conducted by the U.S. Census Bureau.

The actual data is made freely available from the UC-Irvine Machine Learning Repository, and consists of 46 variables (see the appendix for a full list), with features such as race, capital gains, country of birth of parents, and weeks worked in the year. Because the census data is completed by participants on a voluntary basis, some variables have more missing observations than observed observations. This is a challenging statistical problem because a household's decision to respond or not to a given question may be correlated with their household income, our choice variable of prediction. This means that normal regression methods will result in inconsistent estimators. For example, for a standard regression model $Y = X\beta$, we would have:

$$E(y_i|x_i) = x_i\beta + E(y_i|x_j = NA) \tag{1}$$

where x_j represents the variables $j = \{j_1, j_2, \dots, j_k\}$ with x_j not reported. Put intuitively, the second term $E(y_i|x_j = NA)$ describes how we would alter our prediction given that we have NA in columns j . It is possible that wealthier households are more likely to not self-report capital gains, and therefore, we should raise our expected probability of a household having annual income greater than \$50,000 if we observe a missing value. (Another problem all-together is that households could have purposefully mis-reported-a complexity that we avoid all-together to simplify our analysis). The KNN method (described later) seeks to alleviate, but does not solve this problem. Perhaps due to this interaction between household income and missing data, several variables that intuitively would be a very good measure of household income, like wage/hr, weeks worked per year, and capital gains, are not only missing in many observations, but are not good predictors of household income.

One of the most pertinent characteristic of the data set is that the "over \$50K" variable only takes on a value of true for 6.2% of the observations. As we noted in our original proposal, this means that a majority classifiers—predicting that income is under 50K for all observations, would have been "93.8% accurate". However, this is not a useful predictor, since we do not know if a set of new observations would be representative at all of our current data. In addition, it's important to consider weighting the relative importance of assigning a false positive vs. a false negative: for a business/marketing application of the data, it may be more important to predict all of the observations with income $> \$50K$, than to correctly predict households with income under \$50K. In this sense, the majority classifier is undesirable, as it predicts no observations with income $> \$50K$, regardless of information.

3 Logistic Regression

The task of predicting whether a household self reports that its income is over or under \$50,000 is a binary classification, as our outcome takes a value of either true (income over \$50K) or false (income under \$50K). We decided early on in the project to use logistic regression, for reasons to be explained shortly, and sought to find the best classification method within the subset of logistic regression classifiers.

3.1 Explanation of Logistic Regression model

Logistic regression transforms the response based on the predictors to

$$\pi(X) = \frac{e^{(X\beta)}}{1 + e^{(X\beta)}} \quad (2)$$

a transformation which looks like:




Figure 1: A graph of the logistic function for reference, graphed by the authors using the numpy module of python

3.2 Intuition behind logistic regression

This function can be viewed as a CDF (cumulative density function) and closely mirrors that of a normal distribution, and projects the space of $X\beta$ to the interval $[0, 1]$, allowing us to view the result as the probability that our outcome is 1. Our reasons for selecting a logistic regression model are simple: although it is more complicated than a basic regression mode $Y = X\beta$, logistic regression allows us to predict a binary outcome, while resulting in more interpretable coefficients β than "black box" methods such as neural networks or random forest.

4 Initial analysis

Our initial analysis was done in R, and much work was required to find and adapt the necessary statistical methods for use in python. We used the Akaike Information Criteria (AIC), defined as $AIC = 2k - 2\ln(L)$, where k is the number of the parameters in the model, and L is the maximized value of the likelihood function - in our case, the logistical regression equation defined above. By maximizing AIC, we produced a subset of predictors that AIC evaluated as the "best" predictor sets. We noticed that among single-variable predictive methods, the education variable was one of the most effective, with a false positive error rate of 5% and false negative error rate of .7%. However, these rates are not accurate measures of performance, because we are testing our predictor with the same data we used to train. (Correcting our evaluation of model effectiveness will be discussed in the cross validation section).

5 Filling in missing observations with 1NN

Since we were working with somewhat cleaned census data, and all the joy that entails, one of the bigger issues that we had to deal with was missing data. Logistic regression cannot be run on an incomplete dataset. The method we decided to use to implement filling in missing data was a one nearest neighbor algorithm. The distance function we used to implement this was very simple. We selected a series of columns that were present in all observations, and the distance between any two observations was simply the number of mismatches across the various columns. The use of

this distance function relies on the critical assumption that people sharing commonalities in these fields will tend to also share commonalities in the missing fields. Whether or not this is true is debateable, but the filling in missing data was useful for completeness. For every missing entry, we took its nearest neighbor, and used its column values to fill in the missing values for that entry.

In our analysis of the missing data, we realized that many of the columns that had missing data tended to involve information that does not necessarily apply to every person. For example, the field for the previous state someone lived in is not applicable to people who have lived in one state their entire lives. Due to this, we ended up dropping the majority of the columns with missing data, retaining only the columns containing information on the nationality of a person's biological parents. It would be interesting to see if our 1NN method actually led to meaningful results.

While one could argue that this method is not statistically well-founded, the fact that only two variables we use for potential logistic regressions out of many more are filled in with this method, debate over the matter seems of minimal consequence.

6 Brute forcing optimal variable selection

Normally, due to both computational constraints, as well as an appeal to statistical well-foundedness, variable selection is done by applying "intuition" about the predictive power of the variables or via stepwise methods. Given our access to Amazon AWS, the Elastic Compute Cloud, and Elastic Map Reduce, we decided to also deviate from the typical methods that a statistician with "intuition." This amounted to us running regressions across as many combinations of variables as we could. We wrote scripts to take in a list of variables as input and automatically run a logistic regression using those variables, keeping track of false positive and false negative rates on both the training data itself, as well as the test data. This procedure was implemented in two different ways.

6.1 Locally identifying optimal combinations of two to three variables

In this method of implementation, we ran our script on our personal computers to generate logistic regressions for every combination of two to three variables that was possible in our dataset. The choice of only running up to all possible combinations of three variables was due to the sheer volume of possible combinations available. In this implementation, we ran 6,545 logistic regressions on our dataset.

6.2 Pulling out the big guns with Amazon AWS (or trying to)

In this implementation, the goal was to run all the logistic regressions using combinations of four to five variables using Amazon's Elastic Map Reduce and mrjob. This would be another 300,000 logistic regressions, and assuming that each regression takes about 10 seconds (this is what we observed on our personal computers), this would amount to about 37 days of logistic regressions. Essentially, the goal here was to take advantage of parallel computing to make statisticians pale in fear at the amount of compute time we spent on a single modeling problem.

To make this possible, we used the pickle package to package our pandas dataframes, and then loaded them into Amazon's S3 storage server. The development of this implementation was greatly hindered by our usage of the pandas, numpy, and statsmodels packages in Python. As these packages were not native to the EC2 machines, our initial roadblock was simply getting the packages installed and working on the EC2 computers via mrjob.

Our final version of this implementation takes in a text file where each row is a list of variables to perform a single logistic regression on, and via mrjob, performs logistic regression on the variables.

Although we had great plans for this implementation, we discovered that it could only take small lists of variables. Once the number of requested logistic regressions got too large, a map job would inevitably fail, causing the entire job to fail, and forcing us to rerun the simulation.

7 Results

Because of difficulties with mr job on AWS EC2, we took our existing mr job code and deployed it on our personal machines for all 2 and 3-variable combinations. Our results for the best-performing models, as ranked by lowest type 2 (false negative) error for the test data, is as follows:

Variables			Training error		Test error	
			type 1	type 2	type 1	type 2
sex	dividends	capital gains	0.050	0.572	0.050	0.566
dividends	wage		0.079	0.587	0.079	0.584
dividends	capital losses		0.082	0.613	0.082	0.610
dividends	capital losses	wage	0.077	0.628	0.077	0.623

Figure 2: A table of the variables included in top-performing models

8 Cross Validation

An important aspect of our project is that our selected models predict accurately, not only for the sample for which we have data, but for the theoretical population. In this case, our population is less theoretical-it is the U.S. population represented by our census data. Because our prediction models are built by minimizing the error within our sample, it's unreasonable to assume that the model will have as low rates of error as those that are achieved by predicting the data used to generate the model.

This problem is usually addressed by the use of training sets and test sets. The data is initially partitioned, often around 3/4 for a training set and 1/4 for a test set. The training set is used to develop a prediction model, and the resulting model is used on the test set to obtain a theoretically unbiased estimate of model accuracy. The major downside to this approach is that valuable data points are lost when observations are assigned to the test sets, because these observations are no longer used to train the model. In situations where predictive power is important we want to minimize the loss of information that results from assigning data points to the test set.

This is the motivation for cross validation, a procedure in which the data set is divided into k sections, or folds. For each $i = \{1, \dots, k\}$, a model is fit on all data excluding the i th fold, which is used as the test set. By repeating this procedure for all i , we obtain a distribution of the error without losing any information given by assigning observations to a specific test set.

To carry out cross validation, we used (like used for much of our analysis), python with statsmodels and pandas libraries. Several functions (viewable in "nelsoncode.py") were used to divide the data into cross sections, fit each regression model, and report back the error rates. As an input, the function used preloaded data from an S3 bucket, or later, to data stored locally, which was then converted to a pandas data frame.

The below figure shows cross-validated error, divided into false positive and false negative, for a model that uses education to predict household income. The horizontal axis, labeled "k-fold" cv, is the number of folds used in CV. As the number of folds increases, the size of the training set in each iteration increases, and, in general, model accuracy improves. The vertical axis displays the percentage of error.

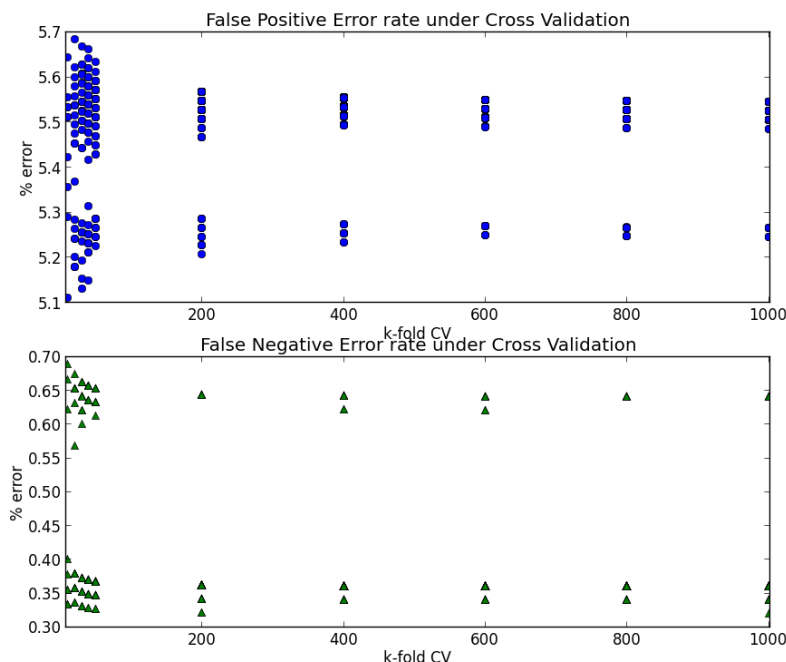


Figure 3: Error rate of education-only model, performed on 5,000 randomly selected observations

A couple of things to note from the results: Although increasing the number of folds seems to decrease the error rate for small k , we do not observe a visible decrease in the error rate for larger k . In addition, we notice a large spread across iterations for the same k -fold CV procedure. This is likely due to the limited variance in our outcome variable-within this 5,000 sample of our data, the number of households with self-reported income over \$50,000 was less than 100. If a couple of difficult-to-predict observations-say, with low education but high income-are placed (by chance) into one specific fold, the error for that fold will be significantly higher than other folds. This is our hypothesis, cannot be examined further given our dataset, and merits further investigation.

Another issue that we wanted to investigate with cross validation was the effect of filling in missing observations using KNN. One member of our team was doubtful that this method would improve the accuracy of our model. To the contrary, filling in observations using data from other observation, that are included in the training set, doesn't seem to add any new information, and might decrease the quality of the model with added variance by making claims on observation with filled-in data.

The below figure shows a similar cross validation procedure, applied to a 5,000-observation sample of the data after having applied the KNN-fill-in-missing-data procedure:




Figure 4: Error rate of education-only model, performed on 5,000 randomly selected observations after filling in missing variables with KNN

The results are extremely interesting for a couple of reasons. We note that compared to the non-filled-in data, the false positive rate is much lower (an average of 4.8% compared to 5.4%) and the false negative rate is slightly higher (.72% compared to .5%). Besides the level of error, the relationship between number of folds and error rates also seems to be different for filled-in data vs. the original data. Looking at the original data, the error rates do not visibly converge. even at 1000-fold CV, there is a .3% range in false positive and false negative error rate between folds. Looking at the cross validation on filled-in data, we see that error rates are more closely clustered, despite the fact that CV was only performed up to 400 folds for the filled-in data, and up to 1000 folds for the original data.

Finally, how might our results of cross-validation change with respect to the size of the data. The previous simulations mentioned were performed on a randomly-selected, 5,000-observation subset of the data. To investigate, we performed the same CV procedure on the entire dataset. This was incredibly computationally-expensive, as the computational complexity of cross validation performed over a number of folds increases exponentially with respect to the number of observations in the data. To ameliorate the computational cost of the simulation, computation was distributed

(manually) among multiple machines, and individual results were combined. The below figure shows the result of the simulation:




Figure 5: Error rate of education-only model, performed on the entire dataset over various fold sizes

Compared to CV on the same model with less observations, the results of the simulation on the entire dataset shows a more narrowly distributed error. The false negative error rate is also slightly higher, although this could be a fluke of the data.

In summery, our cross-validation analysis suggests that the results are likely robust across several datasets, but that we should not expect a complete convergence of error even upon increasing the number of folds.